

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: bmx = pd.read_sas('BMX_I.XPT')
```

```
In [3]: bmx.head()
```

Out[3]:

	SEQN	BMDSTATS	BMXWT	BMIWT	BMXRECUM	BMIRECUM	BMXHEAD	BMIHEAD
0	83732.0	1.0	94.8	NaN	NaN	NaN	NaN	NaN
1	83733.0	1.0	90.4	NaN	NaN	NaN	NaN	NaN
2	83734.0	1.0	83.4	NaN	NaN	NaN	NaN	NaN
3	83735.0	1.0	109.8	NaN	NaN	NaN	NaN	NaN
4	83736.0	3.0	55.2	NaN	NaN	NaN	NaN	NaN

5 rows × 26 columns

In [4]: `bmx.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 9544 entries, 0 to 9543
Data columns (total 26 columns):
SEQN          9544 non-null float64
BMDSTATS      9544 non-null float64
BMXWT         9445 non-null float64
BMIWT         443 non-null float64
BMXRECUM      1073 non-null float64
BMIRECUM      33 non-null float64
BMXHEAD       215 non-null float64
BMIHEAD       0 non-null float64
BMXHT         8769 non-null float64
BMIHT         105 non-null float64
BMXBMI        8756 non-null float64
BMDBMIC       3340 non-null float64
BMXLEG        7110 non-null float64
BMILEG        402 non-null float64
BMXARML       8976 non-null float64
BMIARML       420 non-null float64
BMXARMC       8976 non-null float64
BMIARMC       421 non-null float64
BMXWAIST      8313 non-null float64
BMIWAIST      489 non-null float64
BMXSAD1       6983 non-null float64
BMXSAD2       6983 non-null float64
BMXSAD3       353 non-null float64
BMXSAD4       353 non-null float64
BMDAVSAD      6983 non-null float64
BMDSADCM      446 non-null float64
dtypes: float64(26)
memory usage: 1.9 MB
```

In [5]: `demo = pd.read_sas('DEMO_I.XPT')`

In [6]: `demo.head()`

Out[6]:

	SEQN	SDDSRVYR	RIDSTATR	RIAGENDR	RIDAGEYR	RIDAGEMN	RIDRETH1	RID
0	83732.0	9.0	2.0	1.0	62.0	NaN	3.0	3.0
1	83733.0	9.0	2.0	1.0	53.0	NaN	3.0	3.0
2	83734.0	9.0	2.0	1.0	78.0	NaN	3.0	3.0
3	83735.0	9.0	2.0	2.0	56.0	NaN	3.0	3.0
4	83736.0	9.0	2.0	2.0	42.0	NaN	4.0	4.0

5 rows × 47 columns

In [7]: demo.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 9971 entries, 0 to 9970
Data columns (total 47 columns):
SEQN          9971 non-null float64
SDDSRVYR      9971 non-null float64
RIDSTATR      9971 non-null float64
RIAGENDR      9971 non-null float64
RIDAGEYR      9971 non-null float64
RIDAGEMN      695 non-null float64
RIDRETH1      9971 non-null float64
RIDRETH3      9971 non-null float64
RIDEXMON      9544 non-null float64
RIDEXAGM      4060 non-null float64
DMQMILIZ      6149 non-null float64
DMQADFC       527 non-null float64
DMDBORN4      9971 non-null float64
DMDCITZN      9969 non-null float64
DMDYRSUS      2236 non-null float64
DMDEDUC3      2647 non-null float64
DMDEDUC2      5719 non-null float64
DMDMARTL      5719 non-null float64
RIDEXPRG      1288 non-null float64
SIALANG       9971 non-null float64
SIAPROXY      9970 non-null float64
SIAINTRP      9971 non-null float64
FIALANG       9642 non-null float64
FIAPROXY      9642 non-null float64
FIAINTRP      9642 non-null float64
MIALANG       6977 non-null float64
MIAPROXY      6978 non-null float64
MIAINTRP      6978 non-null float64
AIALANGA      5962 non-null float64
DMDHHSIZ      9971 non-null float64
DMDFMSIZ      9971 non-null float64
DMDHHSZA      9971 non-null float64
DMDHHSZB      9971 non-null float64
DMDHHSZE      9971 non-null float64
DMDHRGND      9971 non-null float64
DMDHRAGE      9971 non-null float64
DMDHRBR4      9575 non-null float64
DMDHREDU      9575 non-null float64
DMDHRMAR      9909 non-null float64
DMDHSEDU      5226 non-null float64
WTINT2YR      9971 non-null float64
WTMEC2YR      9971 non-null float64
SDMVPSU       9971 non-null float64
SDMVSTRA      9971 non-null float64
INDHHIN2      9626 non-null float64
INDFMIN2      9642 non-null float64
INDFMPIR      8919 non-null float64
dtypes: float64(47)
memory usage: 3.6 MB
```

In [8]: merged = bmx.merge(right = demo, on = 'SEQN')

In [9]: *#Question 1: Baby Weights*

In [10]: *#Calculate and display the mean weight of baby boys for each month, from month 0 to 12.
#You'll produce 13 values, one for each month.*

```
In [11]: babies = merged.loc[merged['RIDAGEMN'].isnull() == False]
babies = babies.loc[babies.RIDAGEMN <= 12]
x = babies.loc[babies['DMDHRGND'] == 1,:].groupby(["DMDHRGND", "RIDAGEMN"])["BMXWT"].mean()

x = pd.DataFrame(x)
x1 = x.melt()
print(x)
```

		BMXWT
DMDHRGND	RIDAGEMN	
1.0	5.397605e-79	4.841176
	1.000000e+00	5.694737
	2.000000e+00	6.429412
	3.000000e+00	6.793750
	4.000000e+00	7.650000
	5.000000e+00	8.716667
	6.000000e+00	7.943750
	7.000000e+00	9.150000
	8.000000e+00	8.650000
	9.000000e+00	9.808333
	1.000000e+01	9.858333
	1.100000e+01	9.888889
	1.200000e+01	10.644444

```
In [12]: # Calculate and display the mean weight of baby girls for each month, from month 0 to 12.
y = babies.loc[babies['DMDHRGND'] == 2,:].groupby(["DMDHRGND", "RIDAGEMN"])["BMXWT"].mean()

y = pd.DataFrame(y)
y1 = y.melt()
print(y)
```


		BMXWT
DMDHRGND	RIDAGEMN	
2.0	5.397605e-79	4.469231
	1.000000e+00	5.421429
	2.000000e+00	6.000000
	3.000000e+00	6.687500
	4.000000e+00	7.304762
	5.000000e+00	7.511765
	6.000000e+00	8.278571
	7.000000e+00	8.510526
	8.000000e+00	8.780952
	9.000000e+00	9.585000
	1.000000e+01	9.685714
	1.100000e+01	9.214286
	1.200000e+01	9.936364

In [13]: *#Calculate and display the difference between the mean weights of boys and girls for each month.*
`print(x.values - y.values)`

```
[[ 0.3719457 ]
 [ 0.27330827]
 [ 0.42941176]
 [ 0.10625   ]
 [ 0.3452381 ]
 [ 1.20490196]
 [-0.33482143]
 [ 0.63947368]
 [-0.13095238]
 [ 0.22333333]
 [ 0.17261905]
 [ 0.67460317]
 [ 0.70808081]]
```

In [14]: *#Make a line plot showing two lines: one for boys' mean weights months 0-12, and one for girls' mean weights 0-12 (in a different color). The month will go on the x-axis, and the mean weight will go on the y-axis.*


In [15]: `plt.plot(range(13),y['BMXWT'],'green')
plt.plot(range(13),x['BMXWT'], 'blue')
plt.show()`



In [16]: *#Make a scatterplot with linear regression lines for the baby boys' relationship between age and wt, and the baby girls' relationship between age and wt*

```
In [17]: sns.lmplot(x="RIDAGEMN", y="BMXWT", data=babies, hue = "DMDHRGND")
```

```
Out[17]: <seaborn.axisgrid.FacetGrid at 0x1a175dab9e8>
```



```
In [18]: # Based on our plots and output, we can say that baby boys do weigh more than
          # baby girls on average.
          # First, the subtraction table between the baby boys' average with the baby gi
          # rls' showed up as positive for all the months except two meaning that the aver
          # age weight of baby boys was greater than baby girls for the majority of the ag
          # e.
          # The two months where the baby girls' average was higher only had a slight di
          # fference.
          # Also, the linear regression plot shows a higher slope for baby boys, and the
          # refore, we can conclude that baby boys' average will increase at a more higher
          # rate than baby girls' as they age, meaning the average of baby boys will be
          # higher as age goes up.
          # This is also shown in the plot with the two lines, where the baby boy's aver
          # age is always above the baby girls' average except for two months, month = 6,
          # 8.
```

```
In [19]: #Question 2: Height vs Leg Length vs arm Length

          #For all adults aged 20 and up, who have all three measurements of height, upp
          # er leg length, and upper arm length
          two = merged.loc[merged['RIDAGEYR'] >= 20, :]
          two = two.loc[two['BMXLEG'].isnull() == False]
          two = two.loc[two['BMXHT'].isnull() == False]
          two = two.loc[two['BMXARML'].isnull() == False]
```

```
In [20]: #corr standing height, upper leg
          print(two['BMXHT'].corr(two['BMXLEG']))

          0.7874491670325252
```

```
In [21]: #corr standing height, upper arm  
print(two['BMXARML'].corr(two['BMXHT']))
```

0.797872682253647

```
In [22]: #corr upper leg, upper arm  
print(two['BMXARML'].corr(two['BMXLEG']))
```

0.6293694196844941

```
In [23]: #Make adult age groups by decade: i.e. adults aged 20-29.9, adults aged 30-39.9, ... adults aged  
#70-79.9, adults aged 80+ (7 groups total). (not required, but recommended: use pandas.cut ...  
#this function was not explicitly covered in the notes, but you should be able to read the  
#documentation to learn function usage. If you use this, set option: right=False)  
  
two['range'] = pd.cut(two.RIDAGEYR, [20.00,29.99,39.99,49.99,59.99,69.99,79.99,300], right = False)
```

```
In [24]: #[10 pts] For each age group, calculate the mean of the three values.
a = two.groupby('range')['BMXHT'].mean().reset_index()
b = two.groupby('range')['BMXLEG'].mean().reset_index()
c = two.groupby('range')['BMXARML'].mean().reset_index()

a = pd.DataFrame(a)
b = pd.DataFrame(b)
c = pd.DataFrame(c)

print(a)
print(b)
print(c)
```

	range	BMXHT
0	[20.0, 29.99)	167.879617
1	[29.99, 39.99)	167.488202
2	[39.99, 49.99)	166.563244
3	[49.99, 59.99)	166.351476
4	[59.99, 69.99)	164.977480
5	[69.99, 79.99)	164.728016
6	[79.99, 300.0)	162.097091

	range	BMXLEG
0	[20.0, 29.99)	40.064189
1	[29.99, 39.99)	39.556128
2	[39.99, 49.99)	38.680863
3	[49.99, 59.99)	38.163400
4	[59.99, 69.99)	37.342707
5	[69.99, 79.99)	36.984049
6	[79.99, 300.0)	36.516000


	range	BMXARML
0	[20.0, 29.99)	37.068919
1	[29.99, 39.99)	37.154754
2	[39.99, 49.99)	37.094632
3	[49.99, 59.99)	37.188784
4	[59.99, 69.99)	37.252859
5	[69.99, 79.99)	37.350102
6	[79.99, 300.0)	36.908000

```
In [25]: a['BMXHT'] = a['BMXHT'] - 128
```



```
In [26]: #We want to make a plot with three lines showing the relationship between the
age
#decade and the mean height, leg length, and arm length.

plt.plot(range(7), a['BMXHT'])
plt.plot(range(7), b['BMXLEG'])
plt.plot(range(7), c['BMXARML'])
plt.legend()
plt.show()
```



```
In [27]: #Comment on what you think the data says about the
#guiding question. Include any additional analysis you deem appropriate

#We can see from the mean summary that height and Leg Length decrease with age
whereas arm length stays almost consistent throughout.
#This is probably because as you age, your posture gets worse and therefore th
e Leg length and height decreases.
#We can also see that there is a high correlation between height and both arml
ength and leglength.
#However, the correlation between arm length and leg length is less than the a
bove correlation.
#Looking at the plot we could see height and leg length drastically decreasing
with age, whereas the arm length doesn't change at all.
```

```
In [28]: #Question 3: Education Level and income [25pts]
#Use the variable: INDHHIN2 for household income. Use DMDHREDU for the educati
on level of
#the head of the household.
```

```
In [29]: # Filter to adults aged 20 and older.

adults = merged.loc[merged['DMDHRAGE'] >= 20]
print(adults.shape)

(9487, 72)
```

```
In [30]: #Remove people who are missing, refused to answer, or didn't know the househol
d income or education levels.
```

```
In [31]: adults = adults.loc[adults['INDHHIN2'].isnull() == False]
print(adults.shape)
adults = adults.loc[adults['DMDHREDU'].isnull() == False]
print(adults.shape)
adults = adults.loc[adults['INDHHIN2'] <= 15]
print(adults.shape)
adults = adults.loc[adults['DMDHREDU'] <= 5]

print(adults.shape)
```

```
(9210, 72)
(8974, 72)
(8665, 72)
(8640, 72)
```

```
In [32]: #Remove household income categories 12 ($20,000 and over) and 13 (under $20,000) as they
#don't quite fit in with the other income categories
```

```
In [33]: adults = adults.loc[adults['INDHHIN2'] != 12]
adults = adults.loc[adults['INDHHIN2'] != 13]
```


```
In [34]: # Print the shape of the resulting data before the next step.
print(adults.shape)
```

```
(8262, 72)
```

```
In [35]: # Summarize and display the data to explore the relationship between education
level
#and income. This question is purposely open-ended. You choose how best to summarize and
#display the data to answer the question regarding education level and household income.
```

```
In [36]: sns.boxplot(x="DMDHREDU",y = "INDHHIN2", data = adults)
adults.groupby("DMDHREDU")["INDHHIN2"].mean()
```

```
Out[36]: DMDHREDU
1.0      5.648438
2.0      6.262626
3.0      7.464968
4.0      8.619637
5.0     11.646460
Name: INDHHIN2, dtype: float64
```



```
In [37]: #Looking at the result and looking at the outputted graph, we can see that the
          #mean household income is linearly proportional to educational level.
          #This means that, by average, households with higher education level will result
          #in a higher income level.
          #We see that from education from 1 to 4, the difference of means per education
          #level does not significantly increase the income level as the mean income level
          #increases by about 1.
          #However, from education level 5, the mean income level increases significantly
          #from the previous education level, by about 3 points.
          #This level 5 corresponds to college graduate, and we can therefore make an assumption
          #that college graduates make significantly more money than non-college graduates.
          #Although there is an increasing rate of income by education level, this increasing
          #rate does not seem significant until the last education level = college graduate.
```